

OS examples - Linux Threads

- Thread creation is done through **clone ()** system call
- **clone ()** allows a child task to share the address space of the parent task (process)
 - Flags control behavior

| flag | meaning |
|---------------|------------------------------------|
| CLONE_FS | File-system information is shared. |
| CLONE_VM | The same memory space is shared. |
| CLONE_SIGHAND | Signal handlers are shared. |
| CLONE_FILES | The set of open files is shared. |


```
int clone(int (*fn)(void *), void *child_stack,  
          int flags, void *arg, ...  
          /*pid_t *ptid, struct user_desc *tls, pid_t *ctid*/ );
```

Consumer-Producer problem – Shared Bounded-Buffer (revisited)


- Shared data:

```
#define BUF_SZ 10
typedef struct {
    . . .
} item;


item buffer[BUF_SZ];
int in = 0;
int out = 0;
```




- Buffer needs to be administered and used as a FIFO or Queue



Initial state



After a few
writes/reads -
No wrapping OR
Both 'in' and 'out'
wrapped



After a few
writes/reads -
Only the "in"
wrapped but not
the "out" index

Bounded-Buffer – Producer

```
item next_produced;
while (true) {

    /* wait till next in != out */
    while (((in + 1) % BUF_SZ) == out);

    /* produce an item */
    buffer[in] = next_produced;
    in = (in + 1) % BUF_SZ;
}
```

Bounded Buffer – Consumer

```
item next_consumed;
while (true) {

    /* wait till in != out*/
    while (in == out);

    /* consume an item */
    next_consumed = buffer[out];
    out = (out + 1) % BUF_SZ;
}
```

- Solution is correct, but can only use BUFFER_SIZE-1 elements (why?)
- What happens if both need to access the same location concurrently (e.g. a shared variable or a shared counter?)
 - Synchronization will then be needed

5.1 Background

Producer

```
while (true) {  
    /* Wait for room in the shared buffer */  
    while (counter == BUFFER_SIZE) ;  
  
    /* Write an entry to shared buffer */  
    buffer[in] = next_produced;  
  
    /* Update write pointer and counter */  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

Consumer

```
while (true) {  
    /* Wait for an entry in shared buffer */  
    while (counter == 0);  
  
    /* Read an entry from shared buffer */  
    next_consumed = buffer[out];  
  
    /* Update read pointer and counter */  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```

- Solution uses BUFFER_SIZE elements **but is incorrect** due to the race condition on the shared “counter” variable.


Race Conditions

e.g. if #counter = 0xF230

```
int counter=0;
void thread1()
{
    .
    .

    counter++;

    .
    .
}
```



Memory

```
void thread2()
{
    .
    .

    counter--;

    .
    .
}
```

Race Conditions

- **counter++** could be implemented in machine code as

```
mov r2, #counter  
  
ld  r1, [r2]      register1 = counter  
inc r1            register1 = register1 + 1  
st  r1, [r2]      counter = register1
```


- **counter--** could be implemented in machine code as

```
mov r3, #counter  
  
ld  r2, [r3]      register2 = counter  
dec r2            register2 = register2 - 1  
st  r2, [r3]      counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

| | | |
|----------------------|----------------------------------|-----------------|
| S0: producer execute | register1 = counter | {register1 = 5} |
| S1: producer execute | register1 = register1 + 1 | {register1 = 6} |
| S2: consumer execute | register2 = counter | {register2 = 5} |
| S3: consumer execute | register2 = register2 - 1 | {register2 = 4} |
| S4: producer execute | counter = register1 | {counter = 6} |
| S5: consumer execute | counter = register2 | {counter = 4} |

e.g. if #counter = 0xF230



Race Conditions – cont.

- A **race condition** or **race hazard** is the behavior of a software (or hardware) system where the output is dependent on the sequence or relative timing of the executing threads.
- A **critical race condition** occurs when the order of operations on shared variables causes them to have unexpected or **erroneous** values.
- A **non-critical race condition** occurs when the order of operations on shared variables does not result in an unexpected or erroneous value.
- Critical race conditions result in invalid execution and bugs. Failure to obey mutual exclusion opens up the possibility of corrupting the shared variables.

Race Conditions - cont.

- A critical race condition occurs when **multiple threads** are performing **non-atomic read-modify-write** concurrently or in parallel.
- It is not necessary for two threads writing to a shared variable concurrently, to result in a critical race condition. A read-modify-write needs to exist to cause a critical race condition.
- Examples of read-modify-write operations:
 - Increment and decrement (e.g. `counter++`, `counter--`)
 - test-and-set
 - compare-and-swap
 - accumulate operations (e.g. `counter+=4`)

Race Conditions – cont.

- Race conditions have a reputation of being [difficult to reproduce and debug](#), since the end result is nondeterministic and depends on the relative timing between interfering threads.
- Problems occurring in production systems can therefore disappear when running in debug mode, when additional logging is added, or when attaching a debugger. Thus, a bug that is due to a race condition is often referred to as a ["Heisenbug"](#).
- Thus, it is better to avoid race conditions in the first place and there is no alternative to proper and careful software design.

5.2 Critical Section Problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code (the section that manipulates shared variables using read-modify-write operations)
 - A Process may be changing common variables, updating a table, writing file, etc
 - To avoid race conditions, when one process is in critical section, no other should be in its critical section
- **Critical section problem** requires the design protocol to solve this.
- Each process must ask permission to enter critical section. This happens in the **entry section**. It may follow critical section with an **exit section**, then **remainder section**

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then:
 - The selection of the processes that will enter the critical section next cannot be postponed indefinitely.
 - Only processes that are in their entry section can participate in the selection.--> **no stalls**
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section (and before its request is granted) --> **no starvation**
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes

Critical-Section Handling in OS kernel

- At any instance of time, multiple kernel-mode processes are running concurrently and two or more processes may share data, and thus they have critical sections (where race conditions are possible).
- Two approaches exist for kernel-mode processes:
 - **Preemptive kernels** – allows preemption of process when running in kernel mode (i.e. kernel thread)
 - More responsive since no process can run for too long (thus also suitable for real-time systems)
 - Possibility of race conditions for shared data.
 - **Non-preemptive kernels** – run till kernel mode is exited, and thus blocks (i.e. yields the CPU voluntarily).
 - If kernel code is designed properly, no kernel thread will spend too long in kernel mode.
 - Essentially free of race conditions in kernel mode

5.4 Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- All H/W solutions described in this section are based on idea of **locking**
 - Protecting critical regions via locks
- **Uniprocessors** – could **disable interrupts**
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems since it requires sending a disable interrupts message to all cores.
 - Operating systems using this approach are not broadly scalable
- Modern machines provide special **atomic hardware instructions**
 - **Atomic** = non-interruptible
 - Either test memory word and set value
 - Or swap contents of two memory words

Solution to Critical-section Problem Using Locks

Process A

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder  
section  
} while (TRUE);
```

Process B

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder  
section  
} while (TRUE);
```

test_and_set Instruction

Definition:

```
bool test_and_set (bool *target)
{
    bool rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically (**it is a single machine instruction**)
it is a single machine instruction
1. Returns the original value of the lock variable (*target)
2. Set the new value of lock variable (*target) to “TRUE”.

Using test_and_set()

- Shared Boolean variable lock, initialized to FALSE
- A possible solution to critical section problem?

```
do {  
    /* Wait till lock is false i.e. not locked, then acquire it */  
    while (test_and_set(&lock));  
  
    /* critical section */  
    . . .  
    /* release the lock at the end (i.e. make it false) */  
    lock = false;  
  
    /* remainder section */  
    . . .  
  
} while (true);
```

fetch_and_add Instruction

Definition:

```
int fetch_and_add (int *target, int inc)
{
    int rv = *target;
    *target = *target + inc;
    return rv;
}
```

1. Executed atomically (**it is a single machine instruction**)
2. Returns the original value of the lock variable (`*target`)
3. Set the new value of (`*target`) to `(*target) + inc`.

compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int rv = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return rv;  
}
```

1. Executed atomically
2. Returns the original value of the lock variable (`*value`)
3. Set the variable “value” the value of the passed parameter “new_value” but only if “*value” == “expected”. That is, the swap takes place only under this condition.

Using compare_and_swap

- Shared integer “lock” initialized to 0;
- A possible solution to critical section problem?

```
do {  
    /* Wait for value to be zero (i.e. lock is released), then acquire lock */  
    while (compare_and_swap(&lock, 0, 1) != 0);  
  
    /* critical section */  
    . . .  
    /* release the lock when done with CS */  
    lock = 0;  
  
    /* remainder section */  
    . . .  
} while (true);
```

Bounded-waiting Mutual Exclusion with test_and_set

- **Previous H/W algorithms didn't satisfy the bounded wait requirement.**
- This algorithm uses common data structures:

```
bool waiting[n];  
bool lock;
```
- The variable Key is not shared
- Proof of mutual exclusion:
 - P_i can enter its critical section only if either `waiting[i] == false` OR `key == false`.
 - The value of `key` can become false only if `test_and_set()` is executed. The first process to execute it will find `key == false`; all others must wait.
 - The variable `waiting[i]` can become false only if another process leaves its critical section; only one `waiting[i]` is set to false, maintaining the mutual-exclusion requirement.

```
do {
```

```
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;
```

```
    /* critical section */
```

```
    ...
```

```
    /* Select next process to run  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
  
    if (j == i)  
        lock = false;  
    else  
        waiting[j] = false;
```

```
    /* remainder section starts below*/
```

```
    ...
```

```
    } while (true);
```

Bounded-waiting Mutual Exclusion with test_and_set

- Proof of progress:
 - Since a process exiting the critical section either sets `lock` to false or sets `waiting[j]` to false. Both allow a process that is waiting to enter its critical section to proceed.
- Proof of bounded wait:
 - When a process leaves its critical section, it scans the array `waiting` in the cyclic ordering $(i + 1, i + 2, \dots, n - 1, 0, \dots, i - 1)$. It designates the first process in this ordering that is in the entry section (`waiting[j] == true`) as the next one to enter the critical section. Any process waiting to enter its critical section will thus do so within $n - 1$ turns.

```
do {
```

```
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;
```

entry section



```
    /* critical section */
```

```
    ...
```

```
    /* Select next process to run  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
  
    if (j == i)  
        lock = false;  
    else  
        waiting[j] = false;
```

exit section



```
    /* remainder section starts below*/
```

```
    ...
```

```
    } while (true);
```

5.5 Mutex Locks

- The OS provides abstraction for the hardware tools previously described, particularly since they require some shared lock variables.
- Simplest is **mutex**.
- Usage: Protect a critical section by first **acquire()** a lock then **release()** the lock
 - Lock = Boolean variable indicating if lock is available or not

```
int main() {  
    do {  
        acquire lock  
        critical section  
        release lock  
        remainder section  
    } while (true);  
}
```

An implementation using `atomic acquire()` and `release()`

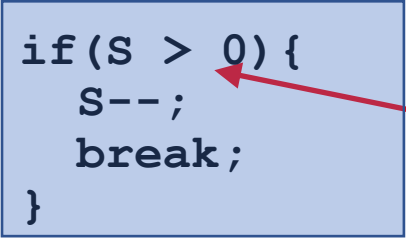
- May be implemented via hardware atomic instructions such as
 - `test_and_set` or `compare_and_swap`
- This lock sometimes referred to as a **spinlock** because it requires **busy waiting**, thus
 - **NOT EFFICIENT.**
 - When used, the critical section must be very short
- A mutex may also be implemented without a spinlock by using **wait queues**. The method is explained in the next section for semaphores.


```
int main() {  
    do {  
        acquire lock  
        critical section  
        release lock  
        remainder section  
    } while (true);  
}
```


5.6 Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore S – integer variable
- **Theoretically**, it can only be accessed via **indivisible (atomic) operations (shown in blue rectangles)** **wait()** and **signal()** (Originally called $P()$ and $V()$)

These are NOT UNIX wait() and signal() API calls

```
wait(S) {  
    while(true) { // busy wait till S>0  
          
    }  
}
```

```
signal(S) {  
      
}
```

$S > 0$ (i.e. 1 and above) indicates that the semaphore is not locked

BLUE RECTANGLES indicate atomic operations

Semaphore Usage

- **Counting semaphore usage** – integer value can range over an unrestricted domain
 - May be used to organize usage of a resource that only allows access to N processes at a time \rightarrow semaphore needs to be initialized to N .
- **Binary semaphore usage** – integer value can range only between 0 and 1
 - Same as a **mutex lock**, except that it has a different polarity (**initialized to 1**)
 - Can synchronization two processes: Consider two processes P_1 and P_2 that require a statement S_1 to happen before S_2

Create a semaphore named “**synch**” and **initialize it to 0**

P1 :

S_1 ;

signal (synch) ;

P2 :

wait (synch) ;

S_2 ;

- **Note that** generally you cannot initialize the semaphore’s value to less than zero. See the man page for `sem_init()`.


Semaphore Implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore **concurrently** (i.e. blue rectangles must be guaranteed to execute **atomically**)
- In similarity to a mutex, processes and threads can be **busy waiting** for the semaphore to become available (i.e. >0)
- A process may thus spend a lot of time in entry sections waiting for the semaphore and not doing any useful work.
- Therefore it may be efficient from the system's point of view to block the process and move it into a waiting queue, and schedule another ready process to run instead.

Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each semaphore has two data items:
 - value (of type integer)
 - pointer to first process in the linked-list queue.
- We define two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```



Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {
```

```
    S->value--;
```

```
    if (S->value < 0) {
```

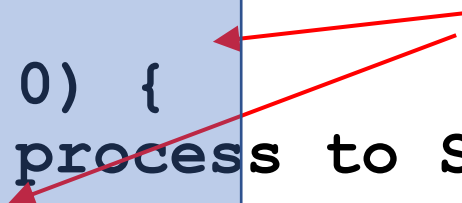
```
        /* add this process to S->list; */
```

```
        block();
```

```
    }
```

```
}
```

Reverse order compared to
that used in busy-waiting



BLUE RECTANGLES indicate
atomic operations

```
signal(semaphore *S) {
```

```
    S->value++;
```

```
    if (S->value <= 0) { /* if someone is in wait queue */
```

```
        /* remove a process P from S->list; */
```

```
        wakeup(P);
```

```
    }
```

```
}
```

Implementation with no Busy waiting (Cont.)

In this implementation, semaphore values are:

- <0 : indicates one or more processes are blocked waiting on the semaphore
 - This is different from previous implementation where the value cannot be <0 .
- $=0$: indicates the semaphore is not available but no process is blocked waiting on it.
- >0 : (i.e. 1 and above) indicates the semaphore is available and thus no process is blocked waiting on it.

Unix/Pthreads Synchronization

- Named semaphores use names that start with '/' and are less than 251 characters.

```
sem_open  
sem_post  
sem_wait  
sem_close  
sem_unlink
```

- Unnamed (anonymous) semaphores use shared variables (for processes or threads) of type `sem_t`.

```
sem_init  
sem_post  
sem_wait  
sem_destroy
```

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0
`wait(S) ;`
`wait(Q) ;`
`...`
`signal(S) ;`
`signal(Q) ;`

P_1
`wait(Q) ;`
`wait(S) ;`
`...`
`signal(Q) ;`
`signal(S) ;`

Priority inversion

- **Priority inversion** is a Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Consider having 3 processes L,M and H with low, medium and high priorities respectively, and consider also a resource R that is shared amongst them.
 - If process L acquired R, and then process H requested R, then H will be blocked.
 - If another process M (priority higher than L and is not requesting R) is ready to run, then it may preempt process L (due to the timer tick for example).
 - This indirectly causes priority inversion and it is sometimes problematic.
- Solved via **priority-inheritance protocol**
 - Priority of process L changes to high when H requests the shared resource, and thus won't be preempted by process M.
- This problem occurred on the Mars Pathfinder's robot in 1997 (running a VxWorks real-time OS).