

5.7 Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
 - Bounded-Buffer Problem
 - Readers and Writers Problem
 - Dining-Philosophers Problem

Bounded-Buffer Problem

- ***BUF_SZ*** elements inside the shared buffer
- Semaphore **`num_full_el`** initialized to the value 0 – keeps track of the number of elements that are full.
- Semaphore **`num_empty_el`** initialized to the value n – keeps track of number of elements that are empty.
- Why use semaphores, when the previous approach seemed to work fine (i.e. using the in and out indices without using any synchronization primitives)?

Bounded-Buffer Problem

- ***BUF_SZ*** elements inside the shared buffer
- Semaphore **`num_full_el`** initialized to the value 0 – keeps track of the number of elements that are full.
- Semaphore **`num_empty_el`** initialized to the value n – keeps track of number of elements that are empty.
- Why use semaphores, when the previous approach seemed to work fine (i.e. using the in and out indices without using any synchronization primitives)?
 - Because of the busy-waiting problem in which a process or thread may be spending valuable CPU time doing nothing but waiting in a loop!
 - We may still use the in and out variables to index a particular buffer in the pool.

Bounded Buffer Problem (Cont.)

- The structure of the **producer** process

```
do {  
    /* produce an item in next_produced */  
    ...  
    /* dec empty sem. */  
    wait(num_empty_el);  
    /* write/produce to an entry*/  
    buffer[in] = next_produced;  
    in = (in + 1) % BUF_SZ;  
    /* inc full sem. */  
    signal(num_full_el);  
} while (true);
```

- The structure of the **consumer** process

```
do {  
    /* dec full sem. */  
    wait(num_full_el);  
    /* read/consume an entry */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUF_SZ;  
    /*inc empty sem.*/  
    signal(num_empty_el);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```

Readers-Writers Problem

- A data set (e.g. a database) is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can read and write
- Problem:
 - Allow multiple readers to read at the same time
 - Only one writer can access the shared data at the same time (i.e. no other writers or other readers are allowed)
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
 - Data set
 - Semaphore `rw_mutex` initialized to 1
 - Semaphore `mutex` initialized to 1 (to protect access to `read_count`)
 - Integer `read_count` initialized to 0

Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1) /* only first reader locks rw_mutex */  
        wait(rw_mutex);  
    signal(mutex);  
  
    ...  
    /* reading dataset is performed, protected by rw_mutex */  
    /* either one writer, or multiple readers at a time  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0) /* only last reader unlocks rw_mutex */  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```

Readers-Writers Problem Variations

- **First** variation – no reader kept waiting unless writer has permission to use shared object
- **Second** variation – once writer is ready, it performs the write ASAP (i.e. no readers are allowed to read till after the writer gets and is done with his access)
- Both may have **starvation** leading to even more variations
- In some systems, the kernel provides a reader-writer locks

Dining-Philosophers Problem

- Philosophers spend their lives alternating between **thinking** and **eating**
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both to eat, then release both when done
- In the case of 5 philosophers
 - Shared data analogy:
 - Each philosopher is a thread
 - Bowl of rice (data set)
 - Semaphore **chopstick [5]** initialized to 1



Dining-Philosophers Problem Algorithm

- The structure of Philosopher i :

```
do {  
    wait (chopstick[i] );  
    wait (chopstick[ (i + 1) % 5] );  
  
    // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```


- What is the problem with this algorithm?

Dining-Philosophers Problem Algorithm (Cont.)

- Deadlock handling
 - Allow at most 4 philosophers to be sitting simultaneously at the table of 5 chopsticks.
 - Use an asymmetric solution
 - An odd-numbered philosopher picks up the left chopstick first and then the right chopstick.
 - An even-numbered philosopher picks up the right chopstick first and then the left chopstick.
 - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section) → we may use **monitors** to implement this method.

5.8 Monitors

- The dining philosophers deadlock may be solved using monitors.
- A monitor is a high-level abstraction that provides a convenient and effective mechanism for process synchronization
- A monitor is an **abstract data type** (i.e. an **object**), internal variables only accessible by code within the procedure
- Only one process may be **active** within the monitor at a time.




5.8 Monitors

```
monitor monitor-name
{
  // shared variable declarations
  procedure P1 (...) { ... }

  procedure Pn (...) {.....}

  Initialization code (...) { ... }
}
```

- But not powerful enough to model some synchronization schemes
- Thus we may use condition variables.




Condition Variables

- Condition variables are variables declared inside a monitor:

```
condition x, y;
```

- Two operations are allowed on a condition variable:
 - `x.wait()` – a process that invokes the operation is suspended until `x.signal()`
 - `x.signal()` – resumes one of processes (if any) that invoked `x.wait()`
- Only one thread/process can be active inside the monitor. The process that issues `x.wait()` becomes inactive, thus allowing others to be active inside the monitor.



Condition Variables – cont.

- Contrast this operation with the `signal()` operation associated with semaphores, which always affects the state of the semaphore:
 - If there are no threads/processes that are waiting on the condition variable `x`, then calling `x.signal()` does not affect the condition variable.
 - Calling `signal(my_semaphore)`, however, increments the semaphore whether there is someone waiting on it or not.

Condition Variables Choices

- If process P invokes `x.signal()`, and process Q was suspended in `x.wait()`, what should happen next?
 - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or Q blocks waiting on another condition (i.e. **immediate effect**).
 - **Signal and continue** – P continues till it leaves the monitor or blocks waiting on another condition, and then Q may become active (i.e. **deferred effect**).
 - Both have pros and cons – language implementer can decide

Monitor solution to dining philosophers

```
monitor DiningPhilosophers
```

```
{
```

```
    enum {THINKING, HUNGRY, EATING} state[5] ;
```

```
    condition cond[5];
```

```
    void pickup (int i) {
```

```
        state[i] = HUNGRY;
```

```
        test_and_signal(i); /* if successful: my state becomes EATING + signal  
                             myself which is wasted cause I am not waiting*/
```

```
        if (state[i] != EATING) cond[i].wait(); /* test(i) did not succeed */
```

```
    }
```

```
    void putdown (int i) {
```

```
        state[i] = THINKING;
```

```
        // test left and right neighbors
```

```
        test_and_signal((i + 4) % 5);
```

```
        test_and_signal((i + 1) % 5);
```

```
    }
```

Monitor solution to Dining Philosophers (Cont.)

```
void test_and_signal(int i) {
    if ((state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING;
        cond[i].signal();
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
} /* end of monitor */
```

Monitor solution to dining philosophers (Cont.)

- Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup(i);`

`EAT`

`DiningPhilosophers.putdown(i);`

- No deadlock, but starvation is possible

Monitor Implementation – Monitors without a condition variable

- Variables

```
semaphore mutex; // (initially = 1)
```

- Each monitor procedure ***F*** will be replaced by

```
wait(mutex) ;  
...  
body of F;  
...  
signal(mutex) ;
```

- Mutual exclusion within a monitor is ensured

Monitor Implementation – Monitors with a condition variables

- For each condition variable **x**, we have:

```
semaphore x_sem; // (initially = 0)
int x_count = 0; // num of processes suspended/waiting
                // on condition variable
```

- We are implementing the **signal-and-wait** type condition variables.
 - A signaling process must wait until the signaled process either **leaves** or **blocks** waiting.
 - Thus we need another semaphore:

```
// signaling process uses next to suspend itself
semaphore next; // (initially = 0)
int next_count = 0; // num of processes suspended
                   // due to signal-and-wait
```

Monitor Implementation – Monitors with a condition variables (cont.)

- The operation **`x.wait()`** can be implemented as:

```
x_count++;  
if (next_count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x_sem);  
x_count--;
```

- The operation **`x.signal()`** can be implemented as:

```
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```

Monitor Implementation – Monitors with a condition variables (cont.)

- With condition variables, each procedure *F* will then be replaced by

```
wait(mutex) ;
```

```
...
```

```
body of F
```

```
...
```

```
// Before exiting, signal other threads that are  
// blocked inside the monitor (i.e blocked  
// themselves when using the signal-and-wait for  
// the condition variable)
```

```
if (next count > 0)
```

```
    signal(next) ;
```


```
else
```

```
    signal(mutex) ;
```


6: Memory Management

- Program translation
- Program relocation
- H/W support for program relocation
- Memory partitioning
- Memory swapping
- Virtual memory
 - Memory segmentation
 - Memory paging
- Virtual Memory management
 - Static virtual memory management and demand paging
 - Dynamic virtual memory management

Program translation




Program translation – linking + loading



Example:

```
static int gVar;  
  
int proc_a(int arg){  
    . . .  
    gVar = 7;  
    put_record(gvar)  
}
```

Source file



```
0000    . . .  
    . . .  
0008    [space for gVar]  
    . . .  
0036    Entry    proc_a  
    . . .  
0220    mov R1,#7  
0224    st R1, [0008]  
0228    push R1  
0232    call "put_record"  
    . . .  
0400    External reference table  
    . . .  
0404    "put_record"      0232  
    . . .  
0500    External definition table  
    . . .  
0540    "proc_a"          0036  
0600    Symbol table (optional)  
    . . .  
0799    Last location in the module
```

Relocatable object file

0000	. . .
. . .	
0008	[space for gVar]
. . .	
0036	Entry proc_a
. . .	
0220	mov R1,#7
0224	st R1, [0008]
0228	push R1
0232	call "put_record"
. . .	
0400	External reference table
. . .	
0404	"put_record" 0232
. . .	
0500	External definition table
. . .	
0540	"proc_a" 0036
0600	Symbol table (optional)
. . .	
0799	Last location in the module

Relocatable object file

0000	Other module
. . .	
1008	[space for gVar]
. . .	
1036	Entry proc_a
. . .	
1220	mov R1,#7
1224	St R1, [1008]
1228	Push R1
1232	call 2334
. . .	
1399	End of proc_a
. . .	
. . .	Start of Other module
2334	Entry put_record
. . .	
2670	(optional symbol table)
2999	(last location in module)

Absolute module

NOTE: Relocatable object files and absolute modules **are binary files** (not text) – Above is only a representation of the actual binary files

Program translation - Relocation issues

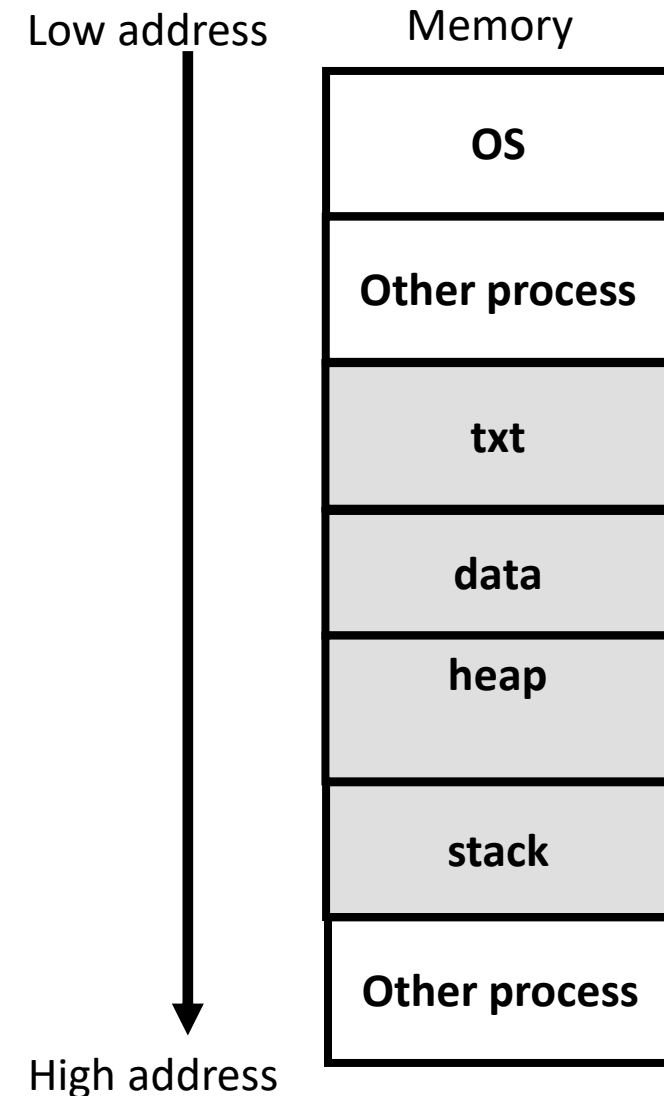
- Referencing static variables:
 - gVar was located at location 0008, but now it is located at 1008. What happens to instructions trying to access gVar ?
 - Clearly, the linker now needs to get to every instruction that accesses gVar and change its reference to location 1008. There are multiple ways of doing this:
 - Linker finds every load or store instruction trying to access the old memory location and changes the address to the new location.
 - Linker uses the optional symbol table, where each relocatable symbol is readily listed and it can easily pin point which instruction is accessing gVar
- Absolute branches
 - **Relative branches** are not a problem since they specify an offset from the current program counter location.
 - Linker needs to modify addresses referenced in **absolute branches** (or calls) just as it did for static variable, e.g. the call to “put_record” gets replaced with the absolute (aka logical) address of “put_record”.

Program translation – file formats

- Source files:
 - Formatted as text files
 - e.g. main.c
- Relocatable object modules and absolute modules:
 - Contain **binary machine instructions** + **information used by linker or loader**
 - Formatted as binary files.
 - Common formats are:
 - COFF: common object file format – commonly used in embedded systems (e.g. microcontrollers: program is loaded into flash memory)
 - PE: portable executable format – used in Windows, extension is .exe
 - ELF: Executable and linkable format – used in Unix/Linux systems.

Program translation – loading

- Memory may be partitioned to accommodate multiple running processes.
- A program is loaded into an area of memory that does not necessarily start at address 0x00.
- Thus more relocation needs to take place at load time.



0000	Other module
. . .	
1008	[space for gVar]
. . .	
1036	Entry proc_a
. . .	
1220	mov R1, #7
1224	st R1, [1008]
1228	push R1
1232	call 2334
. . .	
1399	End of proc_a
. . .	Start of Other module
2334	Entry put_record
. . .	
2670	(optional symbol table)
2999	(last location in module)

Absolute module
e.g. PE (.exe), ELF, COFF, etc.



Other process and OS	
4000	Other modules
. . .	
5008	[space for gVar]
. . .	
5036	Entry proc_a
. . .	
5220	mov R1, #7
5224	st R1, [5008]
5228	push R1
5232	call 6334
. . .	
5399	End of proc_a
. . .	Start of Other module
6334	Entry put_record
. . .	
6999	(last location in module)
Other process	

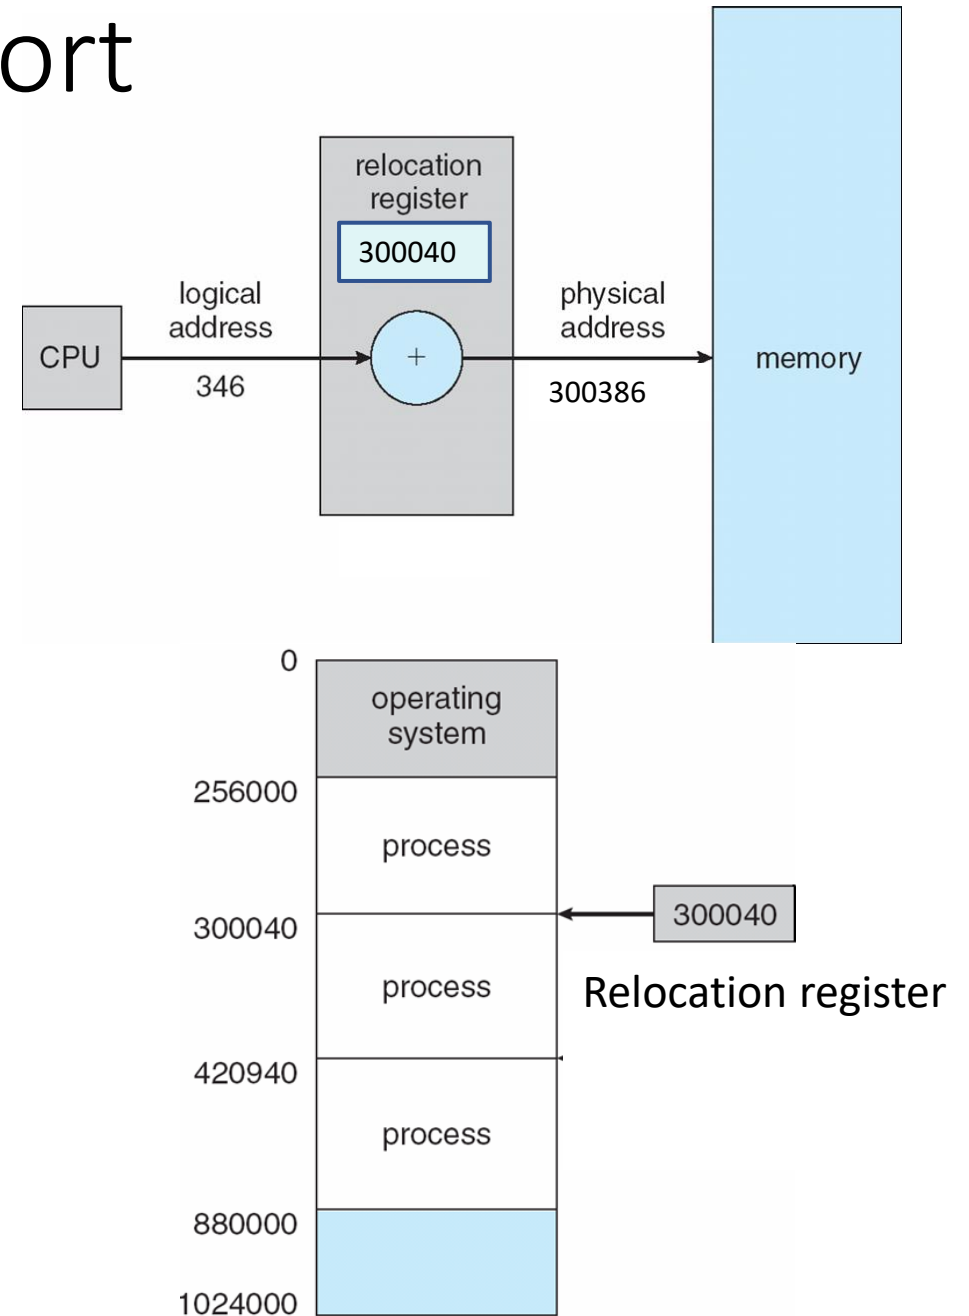
MEMORY

Program translation – cont.

- The process of finding an address in the physical memory (physical address) and attaching it to the relocatable address in the absolute module (= logical address as in your text book) is referred to as **binding**. This binding takes place at **load time**.
- Similarly another binding takes place when attaching addresses in the relocatable object module into addresses in the absolute module. This binding takes place at **compile time**.
- Note that in many systems, it may be that the entire system only runs **one program** and the logical address may be the same as the physical address (e.g. a microcontroller), and thus there may be no need for load-time address binding.
- User programs are designed using the logical address, i.e. the program thinks it is running alone in memory and thinks it starts at address 0.

Hardware relocation support

- A **relocation** register may be used to convert the logical address to a physical address. This register may be also referred to as a **base** register.
- Since a program's logical address starts at 0x0000 but needs to be loaded and executed at a different location in the physical memory, a base register may be summed to every logical address before that address is dispatched to the main memory -> thus it does the job of **binding** instead of the loader.
- Many of the older intel CPUs (e.g. 8088) supported 4 segment registers for: code, data, stack and extra segments.



Hardware Memory protection

- CPU may check every physical memory access generated by a process to be sure it is between **base** and **limit** registers. This hardware thus performs **memory protection**.
- If a process attempts to access an address that is outside its limits, an **exception** takes place and the OS kernel is invoked.
- OS is responsible for setting the correct values in base and limit registers for each process prior to loading it. OS is also in charge of saving that information in the **PCB** during **context switching**.

